

# CSE 221 -- Graduate Operating Systems

---

 [cseweb.ucsd.edu/classes/wi22/cse221-a/hw2.html](https://cseweb.ucsd.edu/classes/wi22/cse221-a/hw2.html)

## CSE 221: Homework 2

---

**Due Thursday, February 24 at 11:59pm**

---

Answer the following questions. For questions asking for short answers, there may not necessarily be a "right" answer, although some answers may be more compelling and/or much easier to justify. But I am interested in your explanation (the "why") as much as the answer itself. Also, do not use shorthand: write your answers using complete sentences.

When grading homeworks, we will grade one question in detail and assign full credit for technical answers to the others.




1. In contrast to blocking locks as provided by semaphores and monitors, a different form of synchronization is known as "non-blocking" or "wait-free" synchronization. With non-blocking synchronization, threads accessing a shared data structure do not block each other. (Linux uses a related kind of synchronization known as read-copy update (RCU) locks, which can have very efficient implementations.)

The general idea is that a shared data structure is accessed via a single shared pointer. Threads just reading the data structure do not acquire a lock and are not blocked by other readers or writers. They happily use the value of that pointer to read data from the data structure. Threads that write (update) the data structure follow these steps:

1. Allocate a separate, private instance of the data structure.
2. Copy the old, shared version of the data structure into the new, private instance (copying just requires reads from the shared version).
3. Update (write to) the new, private instance as needed.
4. Atomically change the pointer referencing the data structure to point to the new instance instead of the old. Any previous reader threads referencing the old data structure are now harmed. Any new reader threads see and see the updates from the new instance. This handles multiple writers, then the last step goes to update the shared pointer. If the pointer is modified since it made a copy, the writer goes back to step 2.

Assuming these

- a. What kinds of data structures (e.g., arrays, linked lists, etc.) work particularly well with this kind of synchronization?
- b. Does this kind of synchronization work for all kinds of data structures?
- c. Can readers starve writers?
- d. Can writers starve readers?
- e. Can readers starve writers?
- f. Can writers starve writers?
- g. Can deadlock occur?
- h. Is it safe for a thread to be arbitrarily suspended or stopped while it is accessing a shared data structure?



2. Exokernel and L4 represent approaches for providing protection and extensibility. Xen represents an approach for providing virtualization and isolation (or, alternately, is an extreme version of extensibility since it goes even beyond Exokernel in exposing the hardware interface to unprivileged code). Consider a Web server as a motivating application-level service running on each of these three system structures, each hosting the OS described in the paper.

For each of the three systems, consider the path a network packet containing an HTTP request takes as it travels from the network interface card to a Web server process running at user level:


- a. Identify the various protection domains in the system for this scenario. Which domains are privileged, and which are unprivileged? (Feel free to draw "boxes-and-kernel-boundary" diagrams if you find them helpful.)

*For example, if the system were standard monolithic Linux, the protection domains would be the kernel and the Web server process with its address space. The kernel is privileged, and the server process unprivileged.*

- b. Describe the path the network packet takes through the protection domains. State the crossing of protection domains.
- c. Argue which service with the network packet is the Web server code into the system and be sure to justify your argument.

- d. Further consider the path the network packet takes through its address space. What happens if there is a page fault? Whose process is the page fault?

*For example, in a standard monolithic Linux, the CPU would raise an interrupt to a Linux kernel. The Linux kernel would allocate a physical page from Linux's free physical page list and update the page table entry with the valid mapping. The Linux kernel would then return from the interrupt.*



3. The notion of "thin" clients has been an attractive idea across the decades. We have seen a few research systems designed to accomodate diskless workstations in the 1980s (V, Sprite, Plan 9). Machines like X Terminals and Sun Rays reached commercial status in the 1990s and 2000s, and Chromebooks perhaps represent the closest thing to a thin client in the modern era. Yet, thin clients have not revolutionized computing. Most of the client devices that we continue to use today (desktops, laptops, smartphones) are full-feature computers still running timesharing operating systems.
- First argue why thin clients were seen as a compelling computing platform, and justify your argument with supporting claims.
  - Then argue why thin clients have not been widely adopted as a platform, also justifying your argument.
  - Which factors in favor of thin clients are still valid today?
- 

